

MATH96012 Project 4

Alexander John Pinches CID: 01201653

December 13, 2019

Question 1.1

As A is a very sparse matrix with its sparsity increasing as n increases creating A as a dense matrix will use lots of memory and be slower for the matrix multiplication as there will be lots of multiplying by zero so instead we store the values and their row and column index this uses less memory for n large and is more efficient when performing the multiplication. Also switching the row and column index allows us to easily multiply by A transposed.

We then only need to perform the multiplication for the values in the sparse form of A with their location in the product determined by their row index.

Question 1.4

The time taken for each method is shown in figure 1. We see that the python implementation is by far the slowest especially for large N with a large jump at 2000. The F90 jacobian implementation is slower than the SGI OpenMP implementation this will be due to the parallelisation which becomes more significant the larger N is. So we should always use the SGI fortran implementation for the fastest results.

Figures

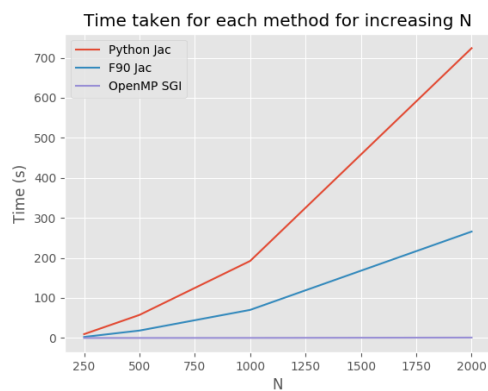


Figure 1: Time taken for increasing N

Question 1.5

We can use the MPE_DECOMP1D subroutine from Argonne National Laboratory to give the start and end index for each process. We can then have each process calculate its portion of Md it will need the start to end columns and rows of A for the multiplication by A and then we can calculate the get the scalars for $e^T e$ and $d^T M d$ for its start and end we then divide these two numbers. To get K we then sum across all processes using `MPI_Allreduce` as we want the sum to be on each process with the operation specified as `MPI_SUM`. Using the k we can then calculate the new x and e for each process between their start and end. We can then calculate $e^T e$ for the new e and get the scalar μ for each process then use `MPI_Allreduce` with `MPI_SUM` like before. Then each process can calculate its new d . We also calculate the max change in x on each process then use `MPI_Allreduce` with `MPI_MAX` to see the max change in x across all processes. Repeating these steps till we reach the max number of iterations or the change in x falls below the tolerance.

The advantage of the MPI approach is we don't have to rely on shared memory so we can distribute the work across computers without having to have memory shared by both. Also, we don't have to wait whilst the process is forked each time we want to do something in parallel. This makes it more scalable as we aren't limited by cores on a single machine and we have less memory overhead from having to copy arrays for each core. However the communication between processes with MPI is much slower as you have to send data from one process to another. Whereas in openmp everything is kept locally in memory which is quickly accessed by each core especially if it's small enough to be kept in the cpus cache.

Question 2.3

As we know that $\max(a_i) \leq \min(nlocal)$ we need to have the $nlocal - 1$ values above and below the last and first values in `ylocal` we store these in `f`. Whose first $nlocal-1$ values are from the process below the next $nlocal$ are calculated by the current process and then the $nlocal - 1$ after this are calculated by the process above. Which we get from the process above and below by them sending their `ylocal` after each update of `ylocal` and storing in `f`. We use `MPI_REDUCE` to collect the sum of $\exp(i\theta)$ for each process and store it in process 0. Taking the absolute value and dividing by N after the loop on process 0. We use `MPI_GATHER` to take `ylocal` from each process and put it into `y` on process 0.

Question 2.4

The best approach would be to split the theta space into a grid. The theta space is a $M \times M$ grid and if we have say N processes we want each process to be given a $\frac{M}{\sqrt{N}} \times \frac{M}{\sqrt{N}}$ space and to send its results to the processes above

below and to the left and right of itself on this grid. This will mean that there will need to be 2 more MPISEND and MPIRECV calls to send the extra the ys calculated by other processes. For a given process with $id = x$ it will need to get the results from and give its results to the process above it in the grid with id $mod((FLOOR(\frac{x}{\sqrt{N}}) - 1) * \sqrt{N} + mod(x, \sqrt{N}), N)$ and below at $mod((FLOOR(\frac{x}{\sqrt{N}}) + 1) * \sqrt{N} + mod(x, \sqrt{N}), N)$ and the process to its left and right on the grid at $(FLOOR(\frac{x}{\sqrt{N}})) * \sqrt{N} + mod(x - 1, \sqrt{N})$ and $(FLOOR(\frac{x}{\sqrt{N}})) * \sqrt{N} + mod(x + 1, \sqrt{N})$. we dont need the results from other processes because of the restrictions on the maximum size of a. We would also of course have to change the sumSin and RHS to allow to sum across an additional axis and to change the dimensions of the arrays to allow for the extra dimension in the 2d case.